

Analyse customer eligibility for home loan on factors such as income, education level, dependents etc.

By - Team Non-Coders

Members - Shivang Jitendra Khare (PG-2), Ishojit Roy Choudhury (PG-2), Sahil Samir (PG-2), Garima Rathi (PGDM - Business Design), Ladi Amarnath Patro (PG-1)

```
In [4]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
```

```
In [5]: train = pd.read_csv("/content/train_ctrUa4K.csv")
train.head()
```

Out[5]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Propo
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	

```
In [6]: test = pd.read_csv("/content/test_1AUu6dG.csv")
test.head()
```

Out[6]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Propo
0	LP001015	Male	Yes	0	Graduate	No	5720	0	110.0	360.0	1.0	
1	LP001022	Male	Yes	1	Graduate	No	3076	1500	126.0	360.0	1.0	
2	LP001031	Male	Yes	2	Graduate	No	5000	1800	208.0	360.0	1.0	
3	LP001035	Male	Yes	2	Graduate	No	2340	2546	100.0	360.0	360.0	Nan
4	LP001051	Male	No	0	Not Graduate	No	3276	0	78.0	360.0	1.0	

```
In [7]: sub = pd.read_csv("/content/sample_submission_49d68Cx.csv")
sub.head()
```

Out[7]:

	Loan_ID	Loan_Status
0	LP001015	N
1	LP001022	N
2	LP001031	N
3	LP001035	N
4	LP001051	N

```
In [8]: train.shape,test.shape
```

Out[8]: ((614, 13), (367, 12))

We have 614 rows and 13 columns in train dataset & 367 rows and 12 columns in test dataset.

Make a copy of the train and test data - if we have to make any changes in these datasets we would not lose the original datasets.

```
In [9]: train_original=train.copy()
test_original=test.copy()
sub_original=sub.copy()
```

Let's Understand Our Data

```
In [10]: train.columns
```

Out[10]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'], dtype='object')

12 Independent variables and 1 Dependent variable.

```
In [11]: train.dtypes
```

```
Out[11]: Loan_ID          object
Gender           object
Married          object
Dependents       object
Education        object
Self_Employed    object
ApplicantIncome  int64
CoapplicantIncome float64
LoanAmount       float64
Loan_Amount_Term float64
Credit_History   float64
Property_Area    object
Loan_Status       object
dtype: object
```

```
In [12]: test.columns
```

```
Out[12]: Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
 'Loan_Amount_Term', 'Credit_History', 'Property_Area'],
 dtype='object')
```

```
In [13]: test.dtypes
```

```
Out[13]: Loan_ID          object
Gender           object
Married          object
Dependents       object
Education        object
Self_Employed    object
ApplicantIncome  int64
CoapplicantIncome float64
LoanAmount       float64
Loan_Amount_Term float64
Credit_History   float64
Property_Area    object
dtype: object
```

We have 3 formats of Data Types - Objects / Int64 / Float64

Univariate Analysis

Independent Variables

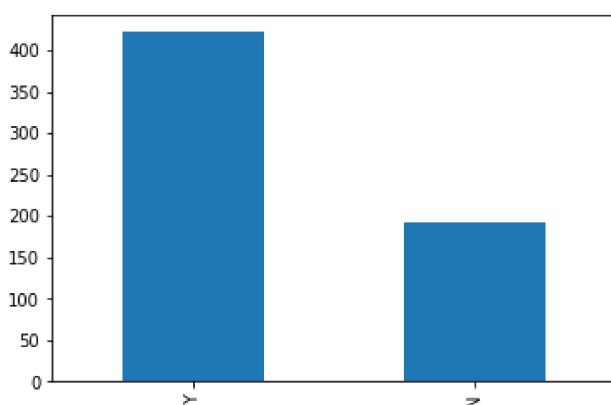
Loan Status Analysis

```
In [14]: train['Loan_Status'].value_counts()
```

```
Out[14]: Y    422
N    192
Name: Loan_Status, dtype: int64
```

```
In [15]: train.Loan_Status.value_counts().plot(kind='bar')
```

```
Out[15]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d982f5c90>
```



Understand the status of Loan.

```
In [16]: train['Loan_Status'].value_counts(normalize=True)
```

```
Out[16]: Y    0.687296
N    0.312704
Name: Loan_Status, dtype: float64
```

Understand loan status in proportional terms instead of numbers. Loan of 68.72% of people was approved.

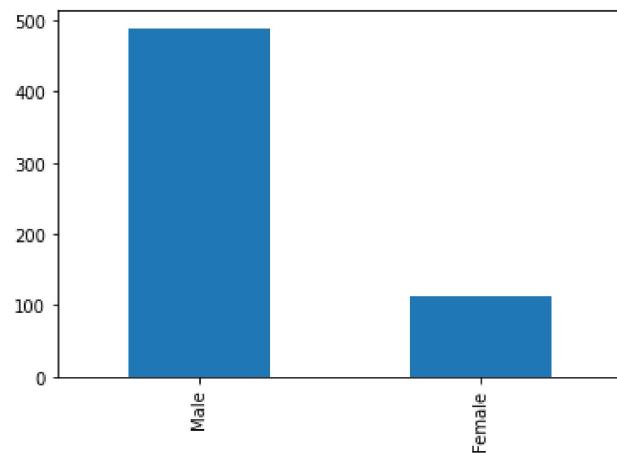
Gender Analysis

```
In [17]: train['Gender'].value_counts()
```

```
Out[17]: Male     489
Female   112
Name: Gender, dtype: int64
```

```
In [18]: train.Gender.value_counts().plot(kind='bar')
```

```
Out[18]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d98209050>
```



```
In [19]: train.Gender.unique()
```

```
Out[19]: array(['Male', 'Female', nan], dtype=object)
```

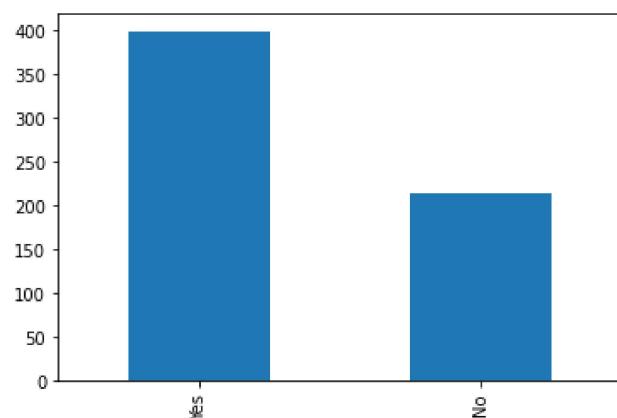
Analysis of Marriage

```
In [20]: train['Married'].value_counts()
```

```
Out[20]: Yes    398  
No     213  
Name: Married, dtype: int64
```

```
In [21]: train.Married.value_counts().plot(kind='bar')
```

```
Out[21]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d96cda510>
```



```
In [22]: train.Married.unique()
```

```
Out[22]: array(['No', 'Yes', nan], dtype=object)
```

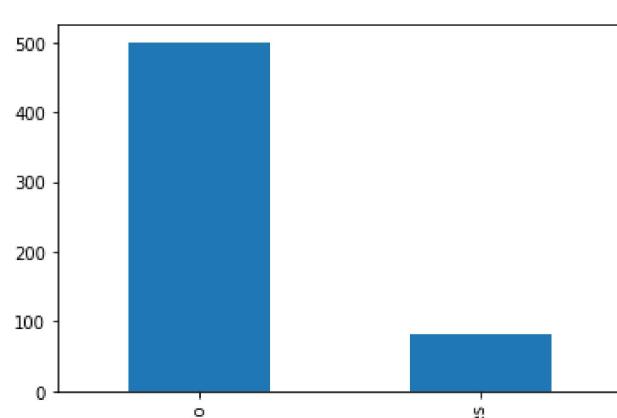
Analysis of Employment

```
In [23]: train['Self_Employed'].value_counts()
```

```
Out[23]: No     500  
Yes      82  
Name: Self_Employed, dtype: int64
```

```
In [24]: train.Self_Employed.value_counts().plot(kind='bar')
```

```
Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d96c59bd0>
```



```
In [25]: train.Self_Employed.unique()
```

```
Out[25]: array(['No', 'Yes', nan], dtype=object)
```

Analysis of Credit History

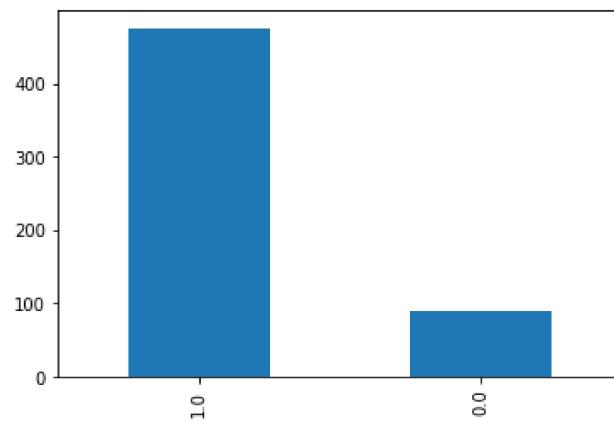
```
In [26]: train['Credit_History'].value_counts()
```

```
Out[26]: 1.0    475  
0.0     89  
Name: Credit_History, dtype: int64
```

```
In [27]: train['Credit_History'].value_counts(normalize=True)
```

```
Out[27]: 1.0    0.842199  
0.0     0.157801  
Name: Credit_History, dtype: float64
```

```
In [28]: train.Credit_History.value_counts().plot(kind='bar')  
Out[28]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d96bcbf90>
```



It can be inferred from above plots that -

1. 81.36% of the applicants are **Males**.
2. 65.13% of the applicants are **Married**.
3. 14.08% of the applicants are **Self Employed**.
4. 84.21% of the applicants have **cleared their doubts**.

Independent Variables

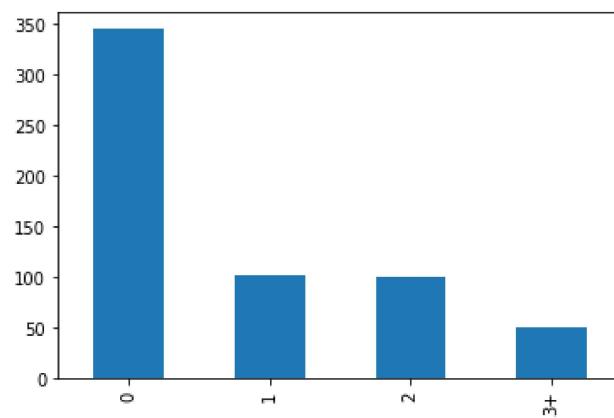
Analysis of Dependents

```
In [29]: train['Dependents'].value_counts()
```

```
Out[29]: 0    345  
1    102  
2    101  
3+    51  
Name: Dependents, dtype: int64
```

```
In [30]: train.Dependents.value_counts().plot(kind='bar')
```

```
Out[30]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d96bb5dd0>
```



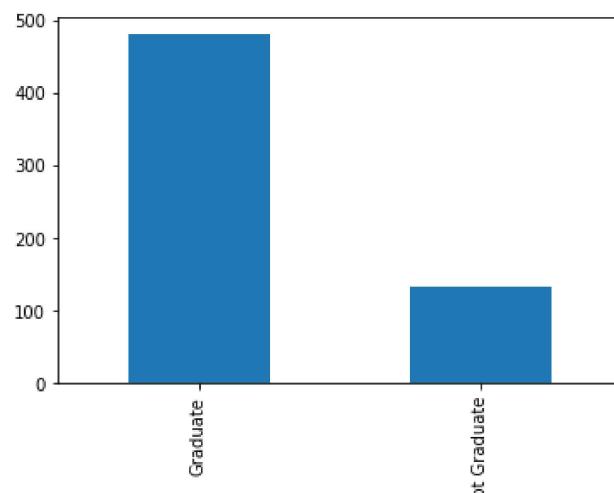
Analysis of Education

```
In [31]: train['Education'].value_counts()
```

```
Out[31]: Graduate      480  
Not Graduate  134  
Name: Education, dtype: int64
```

```
In [32]: train.Education.value_counts().plot(kind= 'bar')
```

```
Out[32]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d96b1d6d0>
```



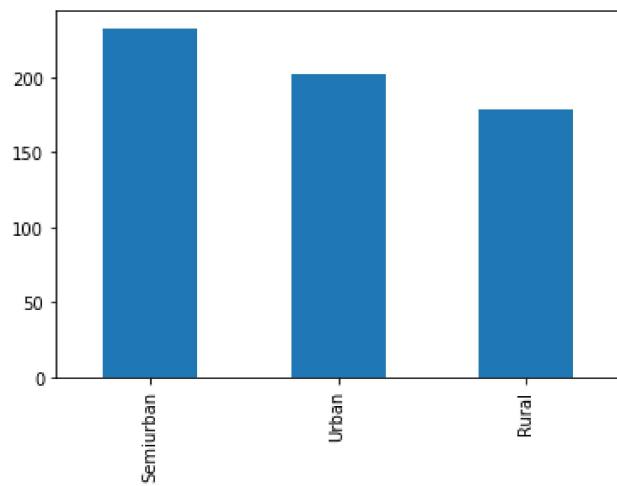
Analysis of Property Area

```
In [33]: train['Property_Area'].value_counts()
```

```
Out[33]: Semiurban   233  
Urban       202  
Rural       179  
Name: Property_Area, dtype: int64
```

```
In [34]: train.Property_Area.value_counts().plot(kind= 'bar')
```

```
Out[34]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d96b309d0>
```



```
In [35]: train['Property_Area'].value_counts(normalize=True)
```

```
Out[35]: Semiurban    0.379479
Urban        0.328990
Rural        0.291531
Name: Property_Area, dtype: float64
```

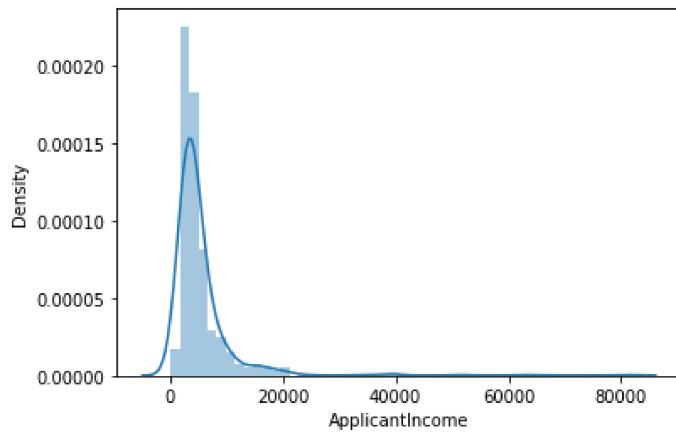
The following inferences can be made from above plots -

1. Most applicants **don't** have dependents.
2. 78.17% of the applicants are **Graduates**.
3. Most applicants are from **Semi-Urban Areas**.

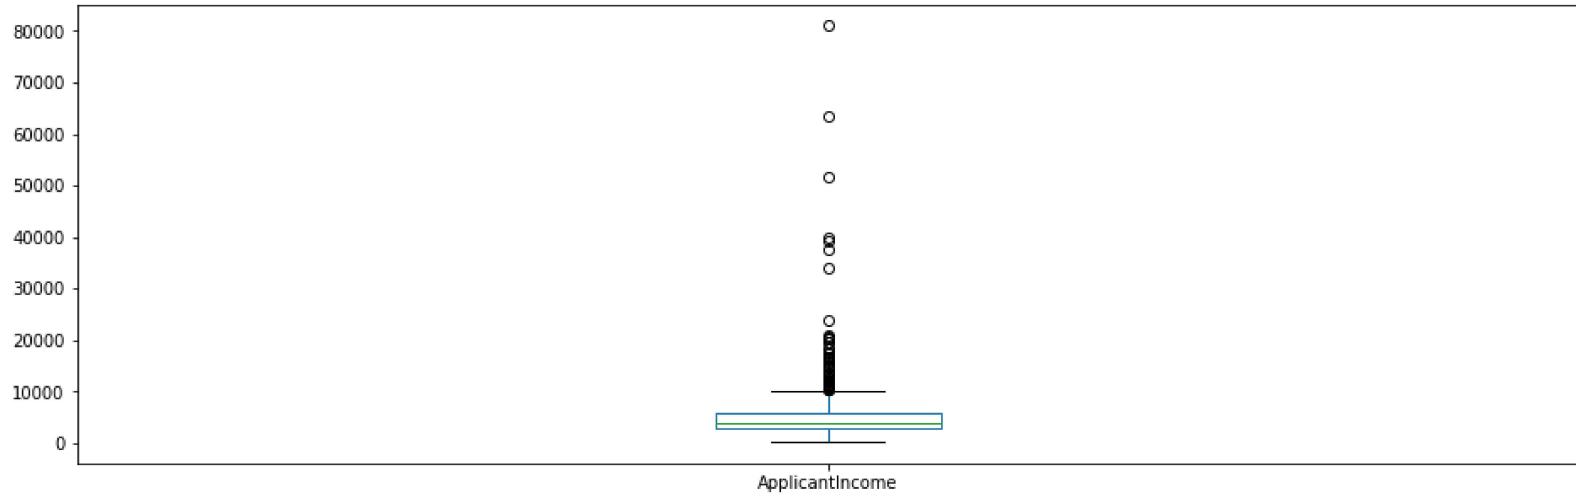
Let's dive deeper into Applicant Income

```
In [71]: sns.distplot(train["ApplicantIncome"])
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



```
In [69]: train["ApplicantIncome"].plot.box(figsize=(16,5))
plt.show()
```



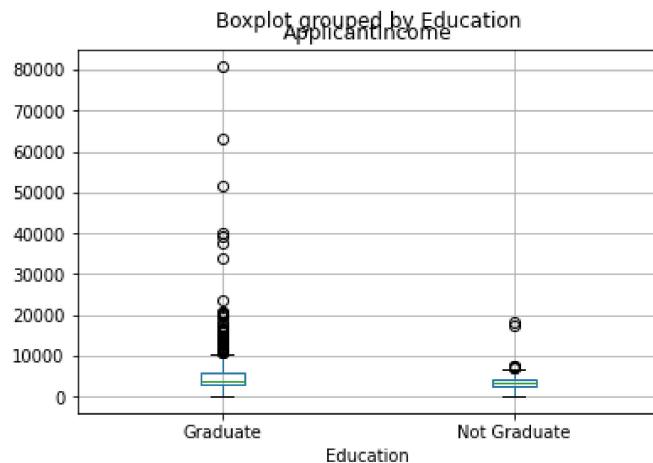
The first plot shows that most of the data in the distribution are towards the left side which means data is **not normally distributed**. The second plot (Box Plot) confirms the presence of a **lot of outliers**. Maybe because the data set contains a lot of people with **different education levels**.

Segregating Applicant Income from Education

```
In [70]: train.boxplot(column='ApplicantIncome', by ='Education')
```

```
/usr/local/lib/python3.7/dist-packages/numpy/core/_asarray.py:83: VisibleDeprecationWarning: Creating an ndarray from ragged nested sequences (which is a list-or-tuple of lists-or-tuples-or ndarrays with different lengths or shapes) is deprecated. If you meant to do this, you must specify 'dtype=object' when creating the ndarray
    return array(a, dtype, copy=False, order=order)
```

```
Out[70]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d969bb710>
```

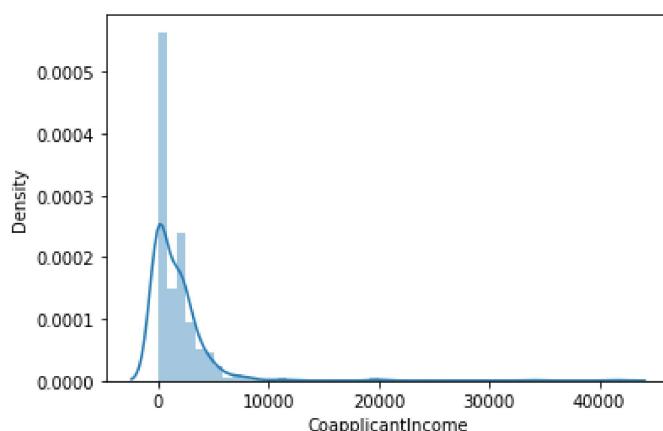


It can be seen that higher number of graduates with higher income appear to be outliers.

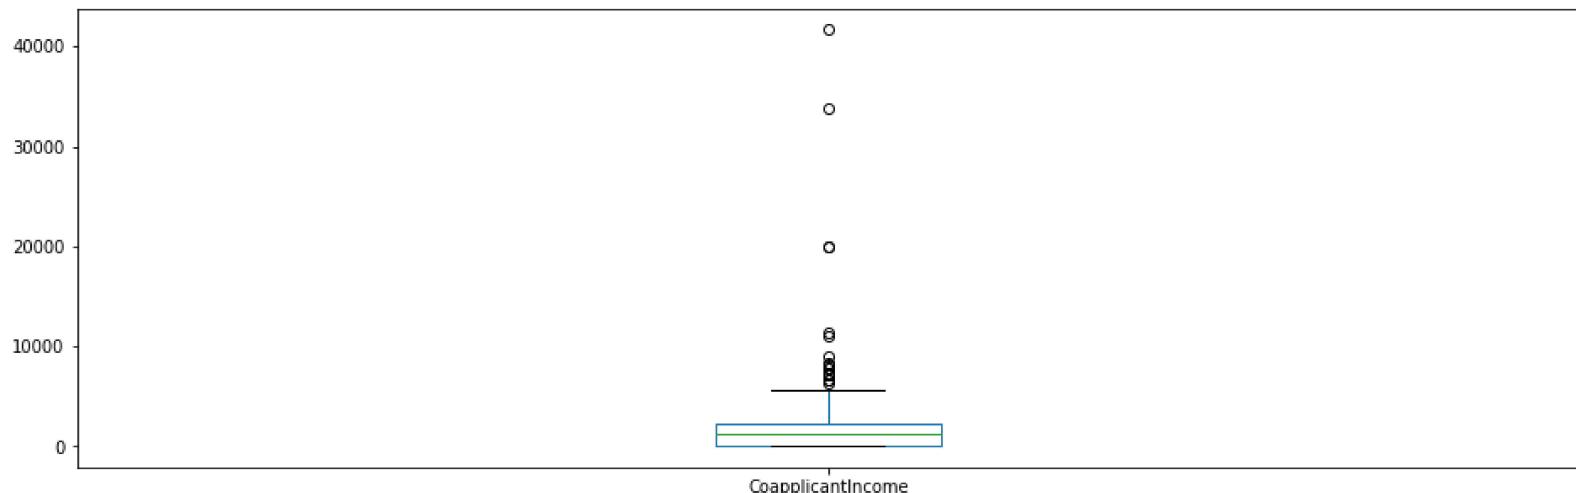
Let's dive deeper into Co-Applicant Data

```
In [39]: sns.distplot(train["CoapplicantIncome"])
plt.show()
```

```
/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
    warnings.warn(msg, FutureWarning)
```



```
In [40]: train["CoapplicantIncome"].plot.box(figsize=(16,5))
plt.show()
```

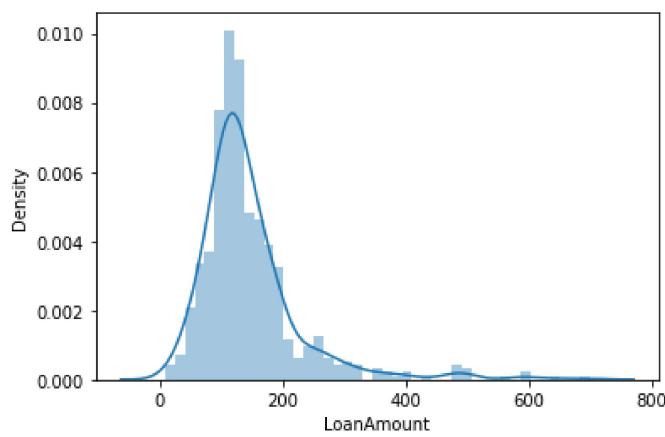


We can see from above plots that the income for most co-applicants range from **0 to 5000**. Also, from 1st plot, we can conclude there are many outliers and data is **not normally distributed**.

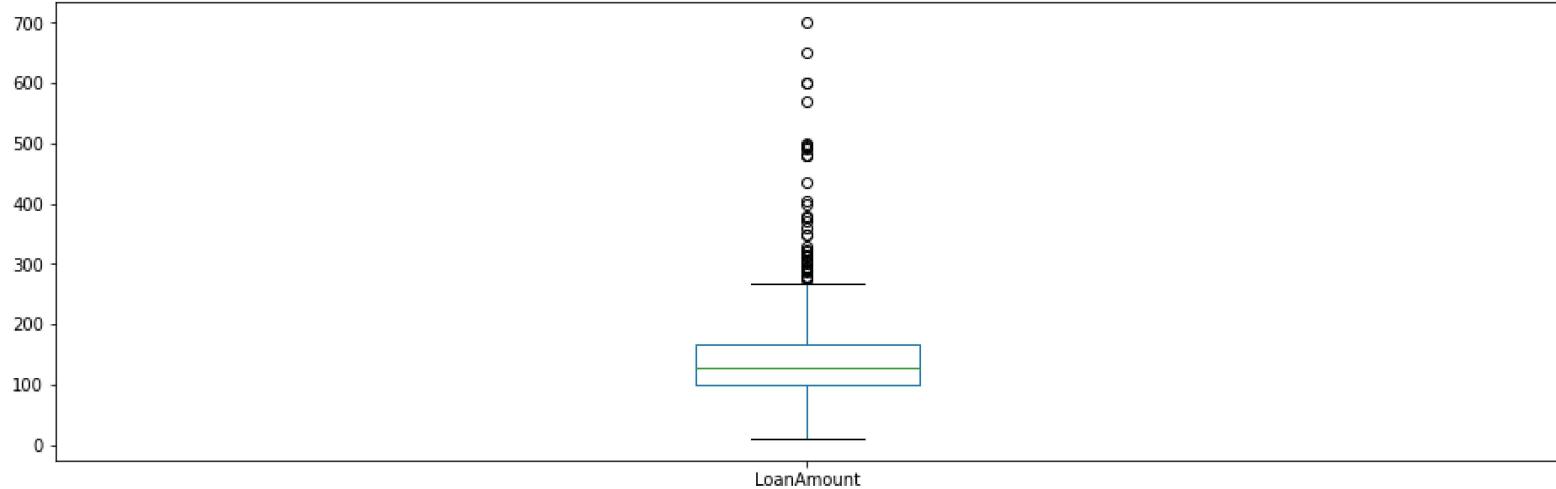
Let's dive deeper into Loan Data

```
In [41]: sns.distplot(train["LoanAmount"])
plt.show()
```

/usr/local/lib/python3.7/dist-packages/seaborn/distributions.py:2619: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
warnings.warn(msg, FutureWarning)



```
In [42]: train["LoanAmount"].plot.box(figsize=(16,5))
plt.show()
```



Distribution is **fairly normal** as per the 1st plot. From 2nd plot it can be seen there are a **lot many outliers**.

```
In [48]: train['Total_Income'] = train['CoapplicantIncome']+train['ApplicantIncome']
train
```

Out[48]:

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History	Pro
0	LP001002	Male	No	0	Graduate	No	5849	0.0	NaN	360.0	1.0	
1	LP001003	Male	Yes	1	Graduate	No	4583	1508.0	128.0	360.0	1.0	
2	LP001005	Male	Yes	0	Graduate	Yes	3000	0.0	66.0	360.0	1.0	
3	LP001006	Male	Yes	0	Not Graduate	No	2583	2358.0	120.0	360.0	1.0	
4	LP001008	Male	No	0	Graduate	No	6000	0.0	141.0	360.0	1.0	
...
609	LP002978	Female	No	0	Graduate	No	2900	0.0	71.0	360.0	1.0	
610	LP002979	Male	Yes	3+	Graduate	No	4106	0.0	40.0	180.0	1.0	
611	LP002983	Male	Yes	1	Graduate	No	8072	240.0	253.0	360.0	1.0	
612	LP002984	Male	Yes	2	Graduate	No	7583	0.0	187.0	360.0	1.0	
613	LP002990	Female	No	0	Graduate	Yes	4583	0.0	133.0	360.0	0.0	

614 rows × 16 columns

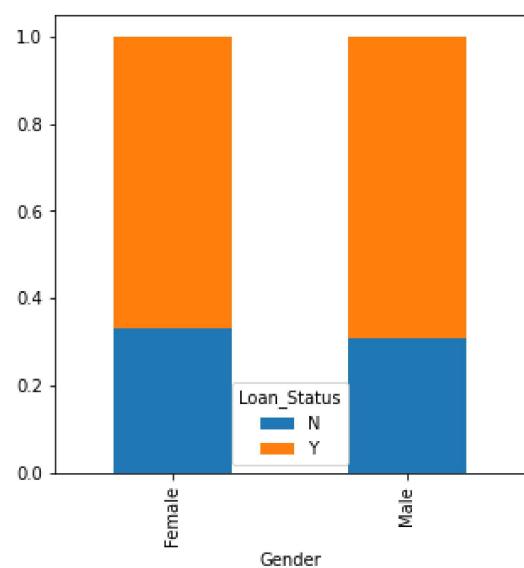
Bivariate Analysis

Hypothesis -

1. Higher income applicants should have higher chances of loan approvals.
2. If loan amount is less, loan approval chances should be high.
3. Lesser the repayment amount, higher the chances to get loan approved.
4. Applicants already repaid previous loan amounts should have higher chance of loan getting approved.

Gender vs Loan Status

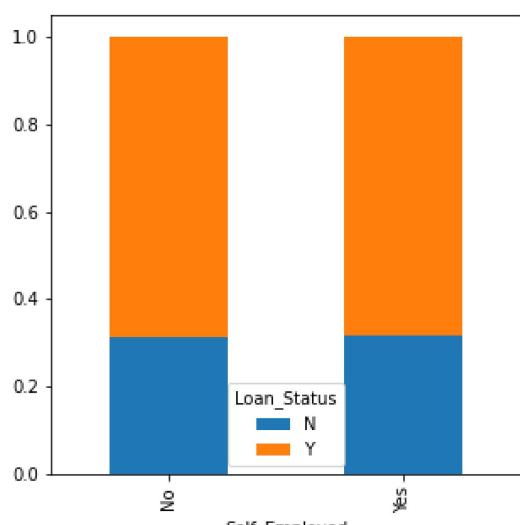
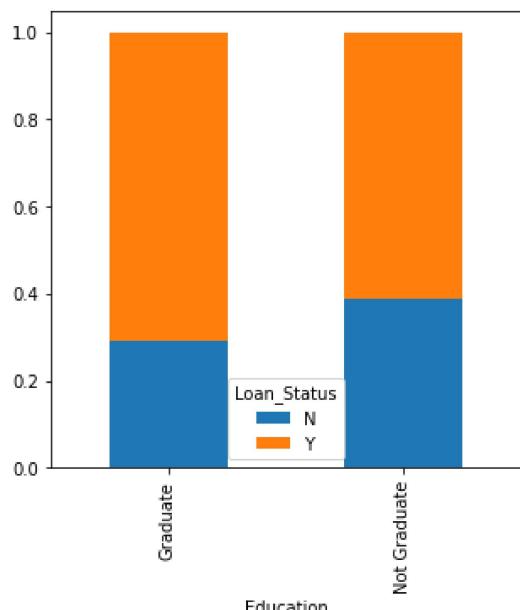
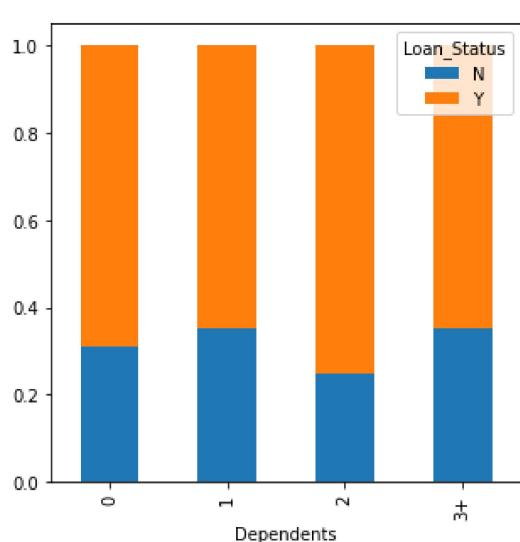
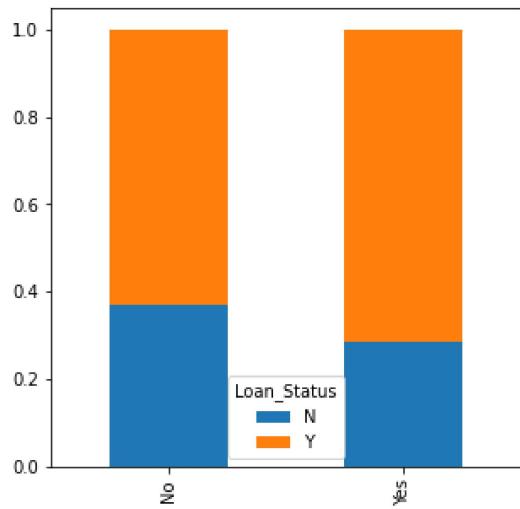
```
In [49]: Gender=pd.crosstab(train["Gender"],train["Loan_Status"])
Gender.div(Gender.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True,figsize=(5,5))
plt.show()
```



Proportion of males and females is almost same for approved and unapproved loans.

Married vs Loan Status / Dependents vs Loan Status / Education vs Loan Status / Self Employed vs Loan Status

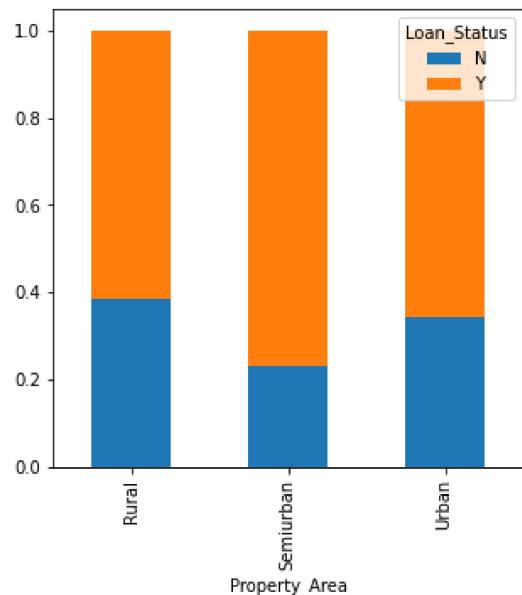
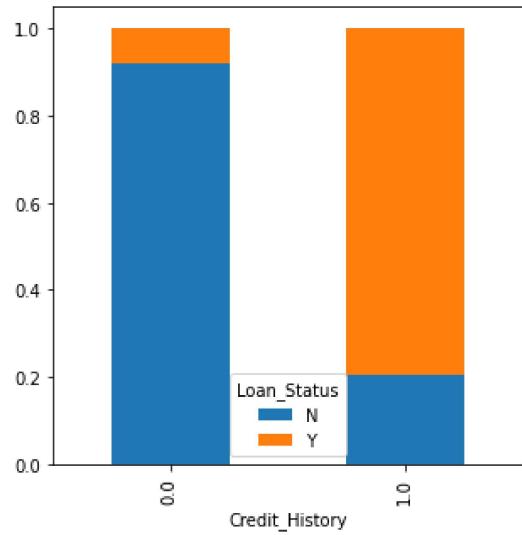
```
In [50]: Married=pd.crosstab(train["Married"],train["Loan_Status"])
Dependents=pd.crosstab(train[ "Dependents"],train[ "Loan_Status"])
Education=pd.crosstab(train[ "Education"],train[ "Loan_Status"])
Self_Employed=pd.crosstab(train[ "Self_Employed"],train[ "Loan_Status"])
Married.div(Married.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True,figsize=(5,5))
plt.show()
Dependents.div(Dependents.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True,figsize=(5,5))
plt.show()
Education.div(Education.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True,figsize=(5,5))
plt.show()
Self_Employed.div(Self_Employed.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True,figsize=(5,5))
plt.show()
```



1. For approved loans, proportion of married applicants is **more**.
2. 1 or 3+ dependents proportion is **same** for approved loans.
3. Nothing significant can be said for **Self-Employed vs Loan Status**.

Credit History vs Loan Status / Property Area vs Loan Status

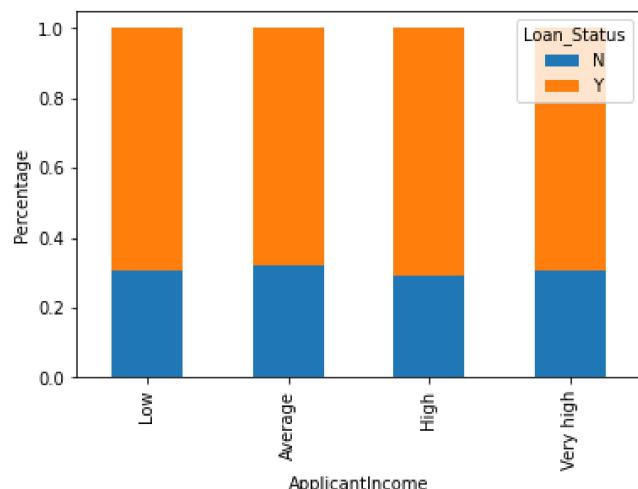
```
In [51]: Credit_History=pd.crosstab(train["Credit_History"],train["Loan_Status"])
Property_Area=pd.crosstab(train["Property_Area"],train["Loan_Status"])
Credit_History.div(Credit_History.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True,figsize=(5,5))
plt.show()
Property_Area.div(Property_Area.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True, figsize=(5,5))
plt.show()
```



1. People with Credit History of 1 have **more** chances of getting loan approved.
2. Semi-urban areas have **higher chance** of getting loan approved than rural or urban areas.

```
In [52]: bins=[0,3000,7000,9000,81000]
group=["Low", "Average", "High", "Very high"]
train["Income_bin"]=pd.cut(train["ApplicantIncome"],bins,labels=group)
Income_bin=pd.crosstab(train["Income_bin"],train["Loan_Status"])
Income_bin.div(Income_bin.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True)
plt.xlabel('ApplicantIncome')
plt.ylabel('Percentage')
```

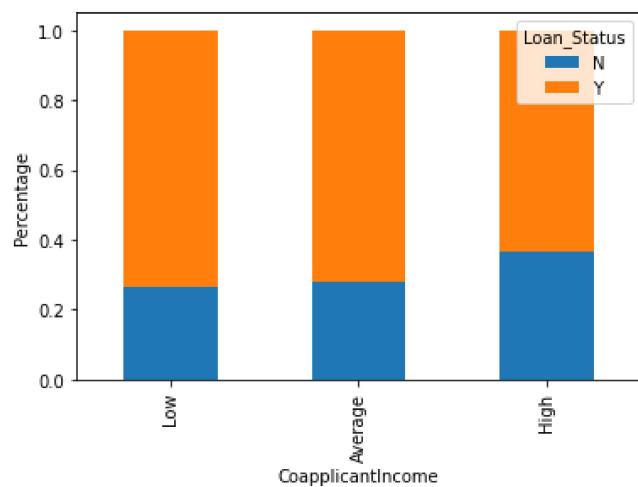
Out[52]: Text(0, 0.5, 'Percentage')



Applicant's income does not affect the chances of loan approvals

```
In [53]: bins=[0,2000,5000,42000]
group=["Low", "Average", "High"]
train["Coapplicant_Income_bin"]=pd.cut(train["CoapplicantIncome"],bins,labels=group)
Coapplicant_Income_bin=pd.crosstab(train["Coapplicant_Income_bin"],train["Loan_Status"])
Coapplicant_Income_bin.div(Coapplicant_Income_bin.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True)
plt.xlabel('CoapplicantIncome')
plt.ylabel('Percentage')
```

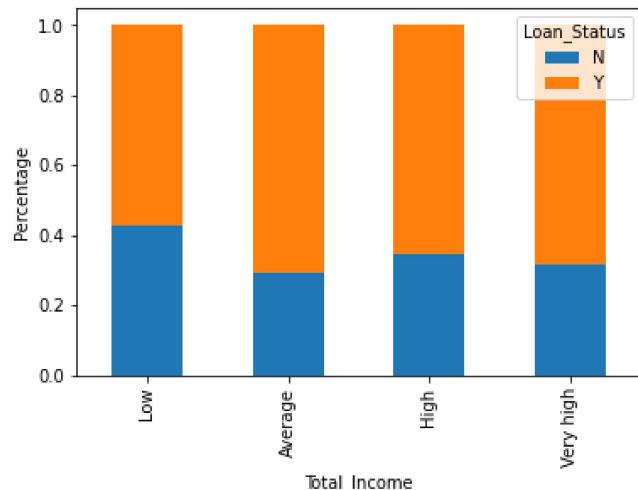
Out[53]: Text(0, 0.5, 'Percentage')



If income of co-applicant is less then Loan approval chances are high but this could be wrong as some applicants may not have a co-applicant at all.

```
In [54]: train["Total_Income"]=train["ApplicantIncome"]+train["CoapplicantIncome"]
bins=[0,3000,7000,9000,81000]
group=["Low", "Average", "High", "Very high"]
train["Total_Income_bin"]=pd.cut(train["Total_Income"],bins,labels=group)
Total_Income_bin=pd.crosstab(train["Total_Income_bin"],train["Loan_Status"])
Total_Income_bin.div(Total_Income_bin.sum(1).astype(float), axis=0).plot(kind='bar',stacked=True)
plt.xlabel('Total_Income')
plt.ylabel('Percentage')
```

Out[54]: Text(0, 0.5, 'Percentage')



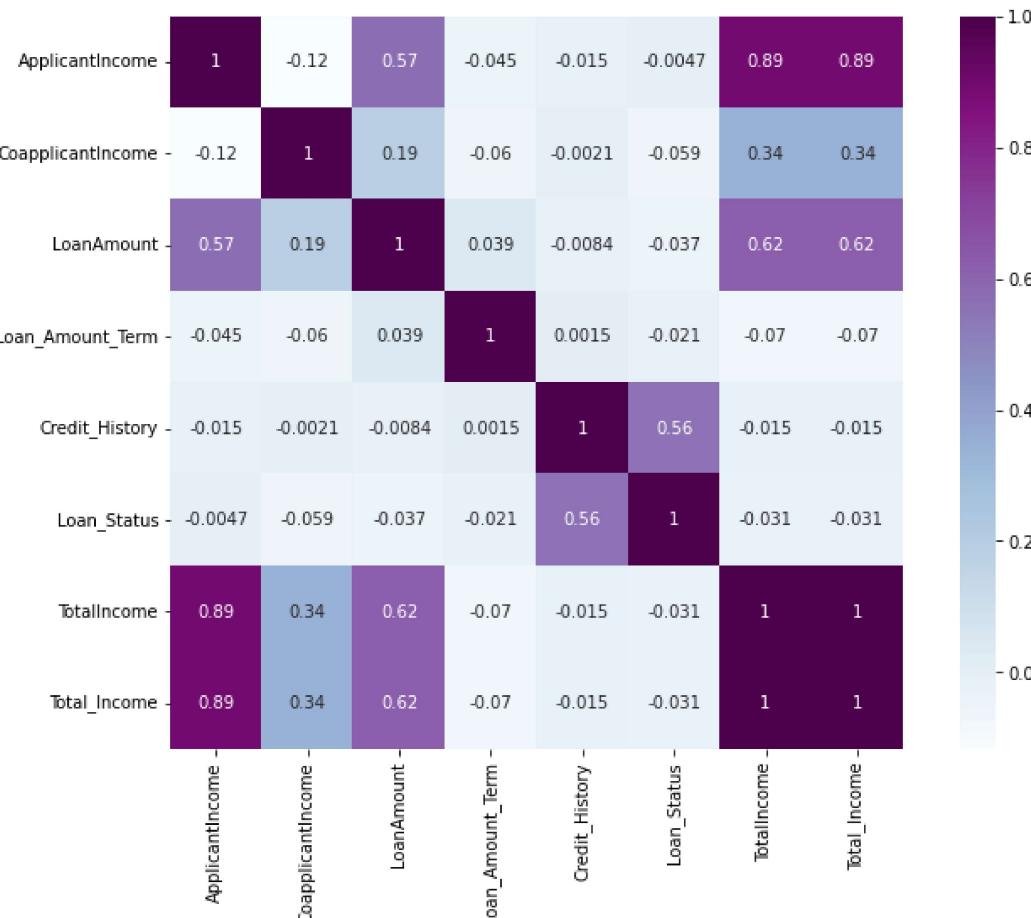
People with low total income have less chances of getting loans approved.

```
In [55]: train=train.drop(['Income_bin', 'Coapplicant_Income_bin'], axis=1)
train["Dependents"].replace('3+', 3,inplace=True)
test["Dependents"].replace('3+', 3,inplace=True)
train["Loan_Status"].replace('N', 0,inplace=True)
train["Loan_Status"].replace('Y', 1,inplace=True)
```

Visualize data via Heat Maps. Variables with darker colors means their correlation is more.

```
In [56]: matrix = train.corr()
f = plt.subplots(figsize=(12,8))
sns.heatmap(matrix,vmax=1,square=True,cmap='BuPu', annot = True)
```

```
Out[56]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4d9435c890>
```



The variables that correlate are (**ApplicantIncome** — **LoanAmount**), (**Credit_History** — **Loan_Status**) and (**LoanAmount** - **CoapplicantIncome**).

Missing Values

```
In [57]: train.isnull().sum()
```

```
Out[57]:
```

Variable	Count
Loan_ID	0
Gender	13
Married	3
Dependents	15
Education	0
Self_Employed	32
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	22
Loan_Amount_Term	14
Credit_History	50
Property_Area	0
Loan_Status	0
TotalIncome	0
Total_Income	0
Total_Income_bin	0

dtype: int64

1. For numerical variables we are using mean and median.
2. For categorical variables we are using mode.

```
In [58]: train["Gender"].fillna(train["Gender"].mode()[0], inplace=True)
train["Married"].fillna(train["Married"].mode()[0], inplace=True)
train["Dependents"].fillna(train["Dependents"].mode()[0], inplace=True)
train["Self_Employed"].fillna(train["Self_Employed"].mode()[0], inplace=True)
train["Credit_History"].fillna(train["Credit_History"].mode()[0], inplace=True)
```

For Loan Amounts Term

```
In [59]: train["Loan_Amount_Term"].value_counts()
```

```
Out[59]:
```

Value	Count
360.0	512
180.0	44
480.0	15
300.0	13
84.0	4
240.0	4
120.0	3
36.0	2
60.0	2
12.0	1

Name: Loan_Amount_Term, dtype: int64

From above, it can be seen that the value of 360 is repeating the most for Loan Amount Term variable. Hence, replacing the missing values with mode of the variable.

```
In [60]: train["Loan_Amount_Term"].fillna(train["Loan_Amount_Term"].mode()[0], inplace=True)
```

For Loan Amounts

We will use median to fill the missing value for loan amounts.

```
In [61]: train["LoanAmount"].fillna(train["LoanAmount"].median(), inplace=True)
```

Now checking the data set

```
In [62]: train.isnull().sum()
```

```
Out[62]:
```

	0
Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0
Property_Area	0
Loan_Status	0
TotalIncome	0
Total_Income	0
Total_Income_bin	0

dtype: int64

All missing values have been filled in the **train** dataset.

For Test Dataset

Similarly, let's fill missing value in test dataset.

```
In [63]:
```

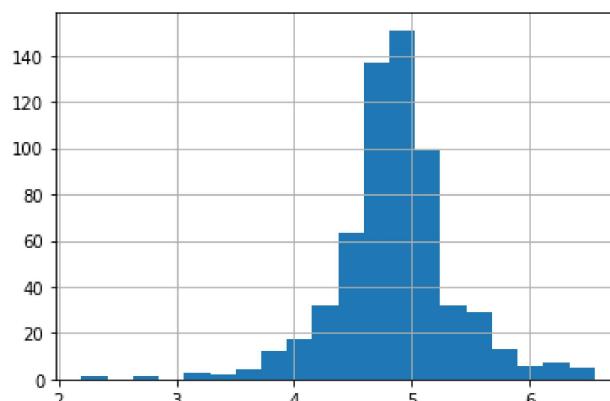
```
test["Gender"].fillna(train["Gender"].mode()[0], inplace=True)
test["Married"].fillna(train["Married"].mode()[0], inplace=True)
test["Dependents"].fillna(train["Dependents"].mode()[0], inplace=True)
test["Self_Employed"].fillna(train["Self_Employed"].mode()[0], inplace=True)
test["Credit_History"].fillna(train["Credit_History"].mode()[0], inplace=True)
test["Loan_Amount_Term"].fillna(train["Loan_Amount_Term"].mode()[0], inplace=True)
test["LoanAmount"].fillna(train["LoanAmount"].median(), inplace=True)
```

Working on Ouliers

Due to the outliers, bulk of the data in the loan amount is at the left and the right tail is longer (as seen in Univariate Analysis). Hence, data is skewed to right. To solve this we are using log transformation method. This will help us get a distribution similar to **Normal Distribution**.

```
In [64]:
```

```
train["LoanAmount_log"] = np.log(train["LoanAmount"])
train["LoanAmount_log"].hist(bins=20)
test["LoanAmount_log"] = np.log(test["LoanAmount"])
```



```
In [65]:
```

```
train["ApplicantIncome_log"] = np.log(train["ApplicantIncome"])
train["ApplicantIncome_log"].hist(bins=5)
test["ApplicantIncome_log"] = np.log(test["ApplicantIncome"])
```

```
/usr/local/lib/python3.7/dist-packages/pandas/core/series.py:726: RuntimeWarning: divide by zero encountered in log
  result = getattr(ufunc, method)(*inputs, **kwargs)
```

